

analysis

October 17, 2021

1 Log Parsing Benchmark Analysis

All log parsing algorithms were run 6 times each on a 12-core Intel Core i7-9750H CPU with 16GB of RAM running Pop_OS! 21.04.

```
[ ]: import os
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import re
```

1.1 Time Analysis

```
[ ]: time_df = pd.DataFrame(columns=["algo", "dataset", "time"])
time_regex = r"(?<=Parsing done. \[Time taken: ]\d:\d+:\d+.\d+"
time_regex_fallback = r"(?<=Parsing done. \[Time: ]\d:\d+:\d+.\d+"

for filename in sorted(os.listdir("outputs")):
    [algo, iteration] = filename.split("_")
    with open("outputs/" + filename, "r") as f:
        times = re.findall(time_regex, f.read())
    if len(times) == 0:
        with open("outputs/" + filename, "r") as f:
            times = re.findall(time_regex_fallback, f.read())
    time_df = time_df.append(
        {
            "algo": algo,
            "dataset": "HDFS",
            "time": pd.Timedelta(times[0]).total_seconds(),
        },
        ignore_index=True,
    )
    time_df = time_df.append(
        {
            "algo": algo,
            "dataset": "BGL",
            "time": pd.Timedelta(times[1]).total_seconds(),
        },
```

```

        ignore_index=True,
    )

time_df

```

```

[ ]:
      algo dataset      time
0      AEL      HDFS  0.362798
1      AEL      BGL  0.270329
2      AEL      HDFS  0.285344
3      AEL      BGL  0.275296
4      AEL      HDFS  0.298853
..      ...      ...      ...
151 Spell      BGL  0.655828
152 Spell      HDFS  0.398777
153 Spell      BGL  0.665933
154 Spell      HDFS  0.391038
155 Spell      BGL  0.660896

```

[156 rows x 3 columns]

```

[ ]: avg_time_df = (
      time_df.groupby(["algo", "dataset"]).mean(numeric_only=False).reset_index()
    )

avg_time_df

```

```

[ ]:
      algo dataset      time
0      AEL      BGL  0.265144
1      AEL      HDFS  0.303317
2      Drain      BGL  0.340670
3      Drain      HDFS  0.375423
4      IPLoM      BGL  0.302273
5      IPLoM      HDFS  0.305821
6      LFA      BGL  0.187106
7      LFA      HDFS  0.210758
8      LKE      BGL  57.855254
9      LKE      HDFS  54.928430
10     Lenma      BGL  2.359167
11     Lenma      HDFS  0.440831
12  LogCluster      BGL  0.203864
13  LogCluster      HDFS  0.196447
14     LogMine      BGL  3.815657
15     LogMine      HDFS  6.311744
16     LogSig      BGL 129.497067
17     LogSig      HDFS  2.398137
18     MoLFI      BGL  26.258781
19     MoLFI      HDFS  3.910462
20     SHISO      BGL  4.869288

```

21	SHISO	HDFS	1.304586
22	SLCT	BGL	1.299598
23	SLCT	HDFS	0.675247
24	Spell	BGL	0.663679
25	Spell	HDFS	0.395011

1.2 F1-measure and Accuracy Analysis

```
[ ]: results_df = pd.DataFrame(columns=["algo", "dataset", "f1_measure", "accuracy"])

for filename in sorted(os.listdir("results")):
    algo = filename.split("_")[0]
    with open("results/" + filename, "r") as f:
        lines = f.readlines()
        lines = [line.strip() for line in lines]
        results_df = results_df.append(
            {
                "algo": algo,
                "dataset": "HDFS",
                "f1_measure": float(lines[1].split(",")[1]),
                "accuracy": float(lines[2].split(",")[1]),
            },
            ignore_index=True,
        )
    results_df = results_df.append(
        {
            "algo": algo,
            "dataset": "BGL",
            "f1_measure": float(lines[1].split(",")[2]),
            "accuracy": float(lines[2].split(",")[2]),
        },
        ignore_index=True,
    )

results_df
```


```
[ ]:
      algo dataset  f1_measure  accuracy
0      AEL   HDFS    0.999984    0.9975
1      AEL   BGL     0.999554    0.9570
2    Drain   HDFS    0.999984    0.9975
3    Drain   BGL     0.999599    0.9625
4    IPLoM   HDFS    1.000000    1.0000
5    IPLoM   BGL     0.999110    0.9390
6      LFA   HDFS    0.999545    0.8850
7      LFA   BGL     0.997902    0.8540
8      LKE   HDFS    1.000000    1.0000
9      LKE   BGL     0.399353    0.1275
```

10	Lenma	HDFS	0.999984	0.9975
11	Lenma	BGL	0.939369	0.6895
12	LogCluster	HDFS	0.951863	0.5460
13	LogCluster	BGL	0.996965	0.8350
14	LogMine	HDFS	0.998840	0.8505
15	LogMine	BGL	0.971268	0.7230
16	LogSig	HDFS	0.991767	0.8495
17	LogSig	BGL	0.934917	0.2265
18	MoLFI	HDFS	0.999984	0.9975
19	MoLFI	BGL	0.999778	0.9660
20	SHISO	HDFS	0.999984	0.9975
21	SHISO	BGL	0.994450	0.7110
22	SLCT	HDFS	0.965812	0.5450
23	SLCT	BGL	0.955247	0.5725
24	Spell	HDFS	1.000000	1.0000
25	Spell	BGL	0.956932	0.7865

```
[ ]: final_df = results_df.merge(avg_time_df, on=["algo", "dataset"])
final_df
```

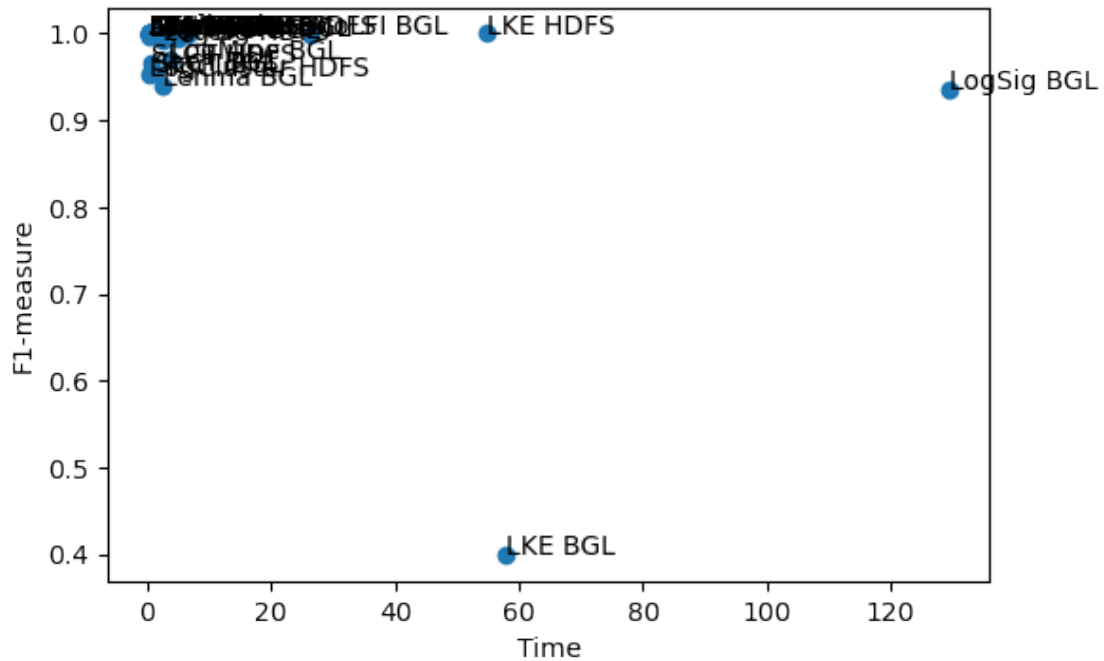
```
[ ]:
      algo dataset  f1_measure  accuracy      time
0      AEL   HDFS    0.999984    0.9975    0.303317
1      AEL   BGL    0.999554    0.9570    0.265144
2    Drain   HDFS    0.999984    0.9975    0.375423
3    Drain   BGL    0.999599    0.9625    0.340670
4    IPLoM   HDFS    1.000000    1.0000    0.305821
5    IPLoM   BGL    0.999110    0.9390    0.302273
6      LFA   HDFS    0.999545    0.8850    0.210758
7      LFA   BGL    0.997902    0.8540    0.187106
8      LKE   HDFS    1.000000    1.0000   54.928430
9      LKE   BGL    0.399353    0.1275   57.855254
10     Lenma   HDFS    0.999984    0.9975    0.440831
11     Lenma   BGL    0.939369    0.6895    2.359167
12 LogCluster   HDFS    0.951863    0.5460    0.196447
13 LogCluster   BGL    0.996965    0.8350    0.203864
14   LogMine   HDFS    0.998840    0.8505    6.311744
15   LogMine   BGL    0.971268    0.7230    3.815657
16   LogSig   HDFS    0.991767    0.8495    2.398137
17   LogSig   BGL    0.934917    0.2265   129.497067
18   MoLFI   HDFS    0.999984    0.9975    3.910462
19   MoLFI   BGL    0.999778    0.9660   26.258781
20   SHISO   HDFS    0.999984    0.9975    1.304586
21   SHISO   BGL    0.994450    0.7110    4.869288
22    SLCT   HDFS    0.965812    0.5450    0.675247
23    SLCT   BGL    0.955247    0.5725    1.299598
24    Spell   HDFS    1.000000    1.0000    0.395011
25    Spell   BGL    0.956932    0.7865    0.663679
```

```
[ ]: plt.figure(dpi=100)
plt.scatter(final_df["time"], final_df["accuracy"])
plt.xlabel("Time")
plt.ylabel("Accuracy")
for i, row in final_df.iterrows():
    plt.annotate(f"{row['algo']} {row['dataset']}", (row["time"],
    ↪row["accuracy"]))
```



Most algorithms are fast, although they have varying accuracies. LKE and LogSig are clear outliers, due to a longer execution time.

```
[ ]: plt.figure(dpi=100)
plt.scatter(final_df["time"], final_df["f1_measure"])
plt.xlabel("Time")
plt.ylabel("F1-measure")
for i, row in final_df.iterrows():
    plt.annotate(f"{row['algo']} {row['dataset']}", (row["time"],
    ↪row["f1_measure"]))
```




We can notice that LKE and LogSig are also outliers in terms of F1-measure over time. These algorithms are not efficient.

Removing them yields the following graph:

```
[ ]: filtered_final_df = final_df[(final_df.algo != "LKE") & (final_df.algo != "LogSig")]


plt.figure(dpi=100)
plt.scatter(filtered_final_df["time"], filtered_final_df["f1_measure"])
plt.xlabel("Time")
plt.ylabel("F1-measure")
for i, row in filtered_final_df.iterrows():
    plt.annotate(f"{row['algo']} {row['dataset']}", (row["time"], row["f1_measure"]))
```



```
[ ]: fig, ax = plt.subplots()


hdfs_df = final_df[final_df.dataset == "HDFS"]
bgl_df = final_df[final_df.dataset == "BGL"]

plt.barh(hdfs_df["algo"], hdfs_df["time"])
ax.invert_yaxis()
ax.set_xlabel("Time (in seconds)")
ax.set_title("Execution Time on HDFS Dataset")
ax.bar_label(ax.containers[0])
fig.set_dpi(100)
plt.show()
```




```
[ ]: fig, ax = plt.subplots()

plt.barh(bgl_df["algo"], bgl_df["time"])
ax.invert_yaxis()
ax.set_xlabel("Time (in seconds)")
ax.set_title("Execution Time on BGL Dataset")
ax.bar_label(ax.containers[0])
fig.set_dpi(100)
plt.show()
```


```
[ ]: fig, ax = plt.subplots()

plt.barh(hdfs_df["algo"], hdfs_df["accuracy"])
ax.invert_yaxis()
ax.set_xlabel("Time (in seconds)")
ax.set_title("Accuracy on HDFS Dataset")
ax.bar_label(ax.containers[0])
fig.set_dpi(100)
plt.show()
```




```
[ ]: fig, ax = plt.subplots()

plt.barh(bgl_df["algo"], bgl_df["accuracy"])
ax.invert_yaxis()
ax.set_xlabel("Time (in seconds)")
ax.set_title("Accuracy on BGL Dataset")
ax.bar_label(ax.containers[0])
fig.set_dpi(100)
plt.show()
```



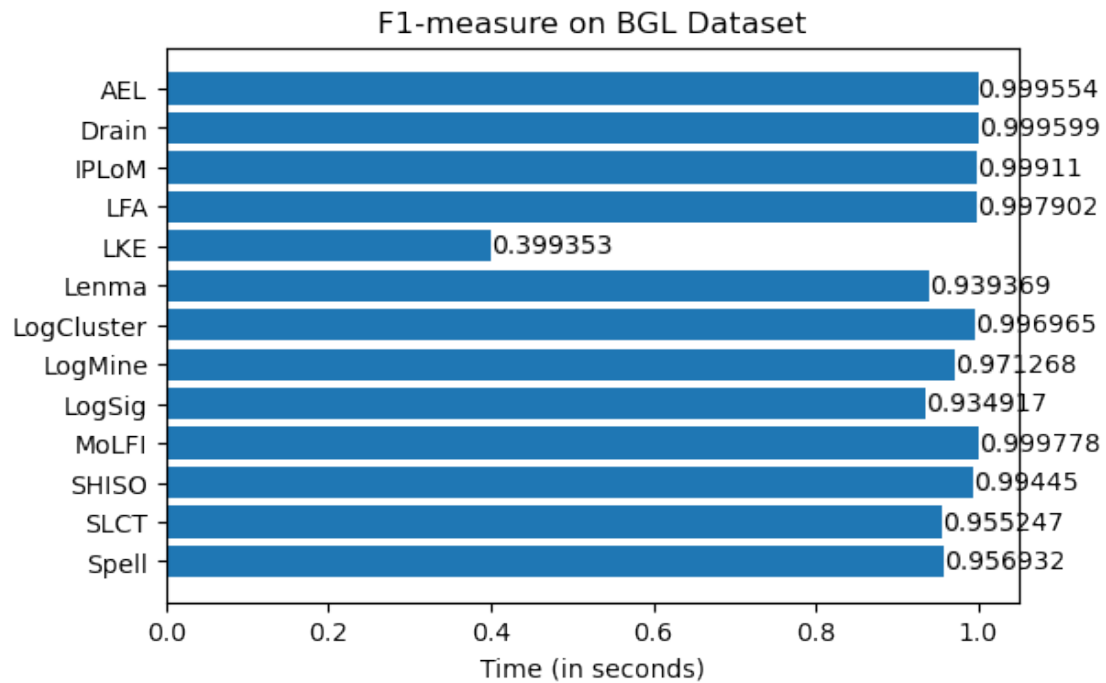
```
[ ]: fig, ax = plt.subplots()

plt.barh(hdfs_df["algo"], hdfs_df["f1_measure"])
ax.invert_yaxis()
ax.set_xlabel("Time (in seconds)")
ax.set_title("F1-measure on HDFS Dataset")
ax.bar_label(ax.containers[0])
fig.set_dpi(100)
plt.show()
```



```
[ ]: fig, ax = plt.subplots()

plt.barh(bgl_df["algo"], bgl_df["f1_measure"])
ax.invert_yaxis()
ax.set_xlabel("Time (in seconds)")
ax.set_title("F1-measure on BGL Dataset")
ax.bar_label(ax.containers[0])
fig.set_dpi(100)
plt.show()
```



```
[ ]: plt.figure(dpi=100)
plt.scatter(final_df["accuracy"], final_df["f1_measure"])
plt.xlabel("Accuracy")
plt.ylabel("F1-measure")
for i, row in final_df.iterrows():
    plt.annotate(
        f"{row['algo']} {row['dataset']}", (row["accuracy"], row["f1_measure"])
    )
```