

analysis

October 17, 2021

1 Log Parsing Benchmark Analysis

All log parsing algorithms were run 6 times each on a 12-core Intel Core i7-9750H CPU with 16GB of RAM running Pop_OS! 21.04.

```
[ ]: import os
import pandas as pd
import matplotlib.pyplot as plt
import re
```

1.1 Time Analysis

```
[ ]: time_df = pd.DataFrame(columns=["algo", "dataset", "time"])
time_regex = r"(?<=Parsing done. \[Time taken: )\d:\d+:\d+.\d+"
time_regex_fallback = r"(?<=Parsing done. \[Time: )\d:\d+:\d+.\d+"

for filename in sorted(os.listdir("outputs")):
    [algo, iteration] = filename.split("_")
    with open("outputs/" + filename, "r") as f:
        times = re.findall(time_regex, f.read())
    if len(times) == 0:
        with open("outputs/" + filename, "r") as f:
            times = re.findall(time_regex_fallback, f.read())
    time_df = time_df.append(
        {
            "algo": algo,
            "dataset": "HDFS",
            "time": pd.Timedelta(times[0]).total_seconds(),
        },
        ignore_index=True,
    )
    time_df = time_df.append(
        {
            "algo": algo,
            "dataset": "BGL",
            "time": pd.Timedelta(times[1]).total_seconds(),
        },
        ignore_index=True,
```

```
)
time_df
```

```
[ ]:      algo dataset      time
0      AEL      HDFS  0.362798
1      AEL      BGL  0.270329
2      AEL      HDFS  0.285344
3      AEL      BGL  0.275296
4      AEL      HDFS  0.298853
..      ...      ...      ...
151 Spell      BGL  0.655828
152 Spell      HDFS  0.398777
153 Spell      BGL  0.665933
154 Spell      HDFS  0.391038
155 Spell      BGL  0.660896
```

[156 rows x 3 columns]

```
[ ]: avg_time_df = time_df.groupby(['algo', 'dataset']).mean(numeric_only=False).
      ↪reset_index()
avg_time_df
```

```
[ ]:      algo dataset      time
0      AEL      BGL  0.265144
1      AEL      HDFS  0.303317
2      Drain      BGL  0.340670
3      Drain      HDFS  0.375423
4      IPLoM      BGL  0.302273
5      IPLoM      HDFS  0.305821
6      LFA      BGL  0.187106
7      LFA      HDFS  0.210758
8      LKE      BGL  57.855254
9      LKE      HDFS  54.928430
10     Lenma      BGL  2.359167
11     Lenma      HDFS  0.440831
12 LogCluster      BGL  0.203864
13 LogCluster      HDFS  0.196447
14 LogMine      BGL  3.815657
15 LogMine      HDFS  6.311744
16 LogSig      BGL  129.497067
17 LogSig      HDFS  2.398137
18 MoLFI      BGL  26.258780
19 MoLFI      HDFS  3.910462
20 SHISO      BGL  4.869288
21 SHISO      HDFS  1.304586
22 SLCT      BGL  1.299598
```

23	SLCT	HDFS	0.675247
24	Spell	BGL	0.663679
25	Spell	HDFS	0.395011

1.2 F1-measure and Accuracy Analysis

```
[ ]: results_df = pd.DataFrame(columns=["algo", "dataset", "f1_measure", "accuracy"])

for filename in sorted(os.listdir("results")):
    algo = filename.split("_")[0]
    with open("results/" + filename, "r") as f:
        lines = f.readlines()
        lines = [line.strip() for line in lines]
        results_df = results_df.append(
            {
                "algo": algo,
                "dataset": "HDFS",
                "f1_measure": float(lines[1].split(",")[1]),
                "accuracy": float(lines[2].split(",")[1]),
            },
            ignore_index=True,
        )
    results_df = results_df.append(
        {
            "algo": algo,
            "dataset": "BGL",
            "f1_measure": float(lines[1].split(",")[2]),
            "accuracy": float(lines[2].split(",")[2]),
        },
        ignore_index=True,
    )

results_df
```

```
[ ]:
      algo dataset  f1_measure  accuracy
0      AEL   HDFS    0.999984    0.9975
1      AEL   BGL     0.999554    0.9570
2    Drain   HDFS    0.999984    0.9975
3    Drain   BGL     0.999599    0.9625
4    IPLoM   HDFS    1.000000    1.0000
5    IPLoM   BGL     0.999110    0.9390
6      LFA   HDFS    0.999545    0.8850
7      LFA   BGL     0.997902    0.8540
8      LKE   HDFS    1.000000    1.0000
9      LKE   BGL     0.399353    0.1275
10    Lenma   HDFS    0.999984    0.9975
11    Lenma   BGL     0.939369    0.6895
```

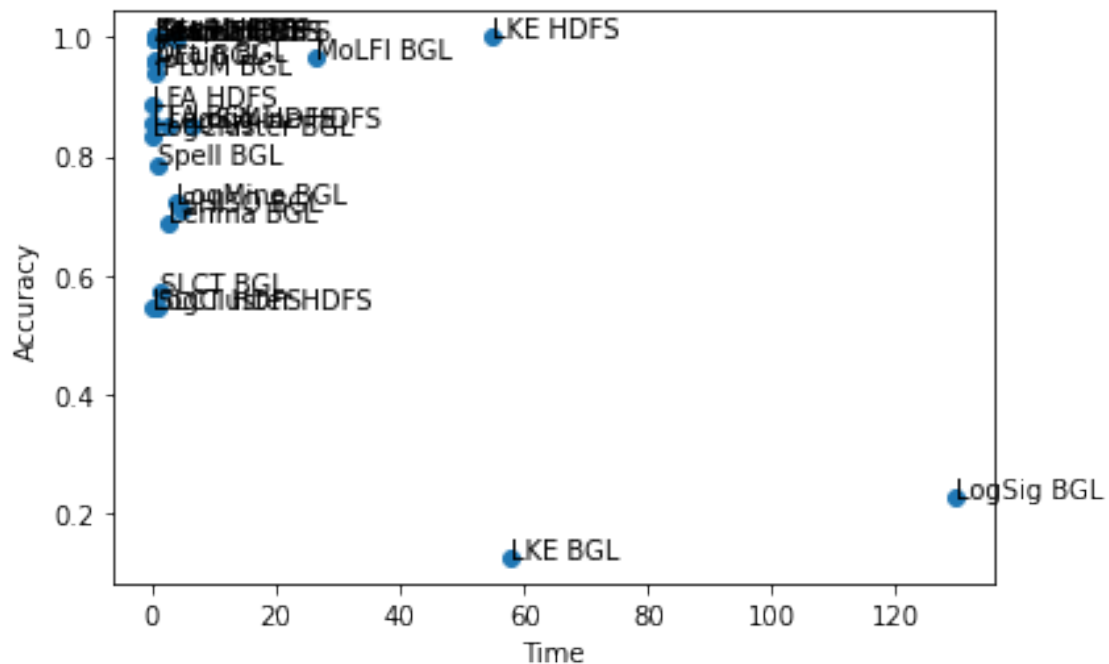
12	LogCluster	HDFS	0.951863	0.5460
13	LogCluster	BGL	0.996965	0.8350
14	LogMine	HDFS	0.998840	0.8505
15	LogMine	BGL	0.971268	0.7230
16	LogSig	HDFS	0.991767	0.8495
17	LogSig	BGL	0.934917	0.2265
18	MoLFI	HDFS	0.999984	0.9975
19	MoLFI	BGL	0.999778	0.9660
20	SHISO	HDFS	0.999984	0.9975
21	SHISO	BGL	0.994450	0.7110
22	SLCT	HDFS	0.965812	0.5450
23	SLCT	BGL	0.955247	0.5725
24	Spell	HDFS	1.000000	1.0000
25	Spell	BGL	0.956932	0.7865

```
[ ]: final_df = results_df.merge(avg_time_df, on=['algo', 'dataset'])
final_df
```

```
[ ]:
```

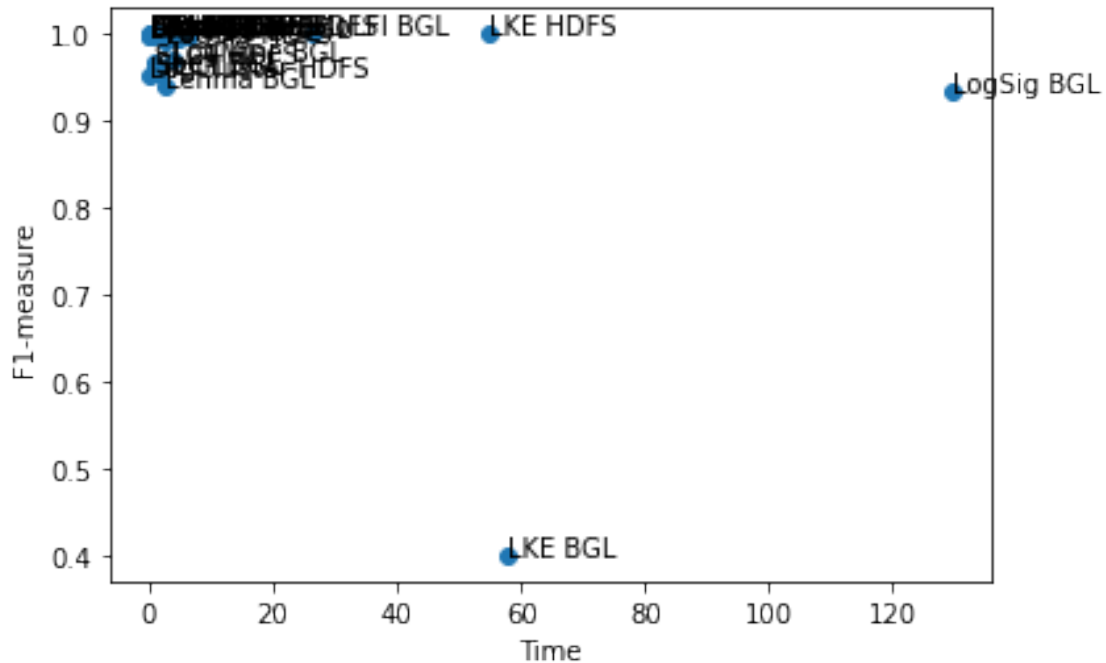
	algo	dataset	f1_measure	accuracy	time
0	AEL	HDFS	0.999984	0.9975	0.303317
1	AEL	BGL	0.999554	0.9570	0.265144
2	Drain	HDFS	0.999984	0.9975	0.375423
3	Drain	BGL	0.999599	0.9625	0.340670
4	IPLoM	HDFS	1.000000	1.0000	0.305821
5	IPLoM	BGL	0.999110	0.9390	0.302273
6	LFA	HDFS	0.999545	0.8850	0.210758
7	LFA	BGL	0.997902	0.8540	0.187106
8	LKE	HDFS	1.000000	1.0000	54.928430
9	LKE	BGL	0.399353	0.1275	57.855254
10	Lenma	HDFS	0.999984	0.9975	0.440831
11	Lenma	BGL	0.939369	0.6895	2.359167
12	LogCluster	HDFS	0.951863	0.5460	0.196447
13	LogCluster	BGL	0.996965	0.8350	0.203864
14	LogMine	HDFS	0.998840	0.8505	6.311744
15	LogMine	BGL	0.971268	0.7230	3.815657
16	LogSig	HDFS	0.991767	0.8495	2.398137
17	LogSig	BGL	0.934917	0.2265	129.497067
18	MoLFI	HDFS	0.999984	0.9975	3.910462
19	MoLFI	BGL	0.999778	0.9660	26.258780
20	SHISO	HDFS	0.999984	0.9975	1.304586
21	SHISO	BGL	0.994450	0.7110	4.869288
22	SLCT	HDFS	0.965812	0.5450	0.675247
23	SLCT	BGL	0.955247	0.5725	1.299598
24	Spell	HDFS	1.000000	1.0000	0.395011
25	Spell	BGL	0.956932	0.7865	0.663679

```
[ ]: plt.scatter(final_df['time'], final_df['accuracy'])
plt.xlabel('Time')
plt.ylabel('Accuracy')
for i, row in final_df.iterrows():
    plt.annotate(f"{row['algo']} {row['dataset']}", (row['time'],
    ↪row['accuracy']))
```



Most algorithms are fast, although they have varying accuracies. LKE and LogSig are clear outliers, due to a longer execution time.

```
[ ]: plt.scatter(final_df['time'], final_df['f1_measure'])
plt.xlabel('Time')
plt.ylabel('F1-measure')
for i, row in final_df.iterrows():
    plt.annotate(f"{row['algo']} {row['dataset']}", (row['time'],
    ↪row['f1_measure']))
```



We can notice that LKE and LogSig are also outliers in terms of F1-measure over time. These algorithms are not efficient.

Removing them yields the following graph:

```
[ ]: final_df = final_df[(final_df.algo != 'LKE') & (final_df.algo != 'LogSig')]
plt.scatter(final_df['time'], final_df['f1_measure'])
plt.xlabel('Time')
plt.ylabel('F1-measure')
for i, row in final_df.iterrows():
    plt.annotate(f"{row['algo']} {row['dataset']}", (row['time'],
    ↪row['f1_measure']))
```

